

IO Enhancements

Jules Kouatchou

NASA GSFC Code 606 (ASTG)
Greenbelt, Maryland 20771
`Jules.Kouatchou@nasa.gov`

October 25, 2018

Agenda

IO Enhancements

- Stream IO
- Asynchronous I/O



Stream I/O - 1

Stream access is a new method for allowing fine-grained, random positioning within a file for read/write operations.

- Complements pre-existing DIRECT and SEQUENTIAL access
- Advantages:
 - Random access (as with DIRECT)
 - Arbitrary record lengths (as with SEQUENTIAL)
- Disadvantages:
 - Presumably poorer performance than both DIRECT and SEQUENTIAL
 - Lack of record separators increases risk of inability to read file under small changes.
 - Index for positioning within file might be less natural than those for DIRECT.



Stream I/O - 2

```
1  OPEN(unit, ACCESS = STREAM)
2  ! both formatted and unformatted files
3
4  READ(unit, POS=n) x,y,z
5  ! File starts at position POS=1
6  ! Position is specified in file storage units -
7  ! usually bytes
8
9  INQUIRE(unit, POS=currentPosition , ...)
10 ! Formatted I/O must use POS obtained from INQUIRE()
11 ! (or POS=1)
```



Examples

Check the files:

- *writeUstream.F90*
- *readUstream.F90*



Asynchronous I/O - 1

Potential performance enhancement allowing some I/O operations to be performed in parallel with other computations.

- To open a file for asynchronous operations, the new optional keyword **ASYNCHRONOUS** is used.
- An asynchronous read/write operation is initiated with the same keyword.
- An optional keyword, **ID**, can be used to return a handle for later use in identifying specific pending operations



Asynchronous I/O - 2

```
1  OPEN(10, ..., ASYNCHRONOUS=yes)
2  WRITE(10,..., id=id, ASYNCHRONOUS=yes) A
3  CALL do_something(...) ! Not involving A here
4  WAIT(10, id=id) ! Blocks here until A has been written
5  CALL do_something(...) ! OK to use A here
```



Asynchronous I/O - 3

If the asynchronous file access is performed in a procedure other than the one called for OPEN, the data involved has to be declared with asynchronous attribute

```
1  OPEN(10, ..., asynchronous=yes)
2  CALL async_write(10, A, id)
3  CALL do_something_else_here()
4  WAIT(10, id=id)
5  ...
6  SUBROUTINE async_write(iu, data, id)
7      INTEGER, INTENT(IN) :: iu
8      INTEGER, INTENT(IN), DIMENSION(:), ASYNCHRONOUS :: data
9      INTEGER, INTENT(OUT) :: id
10     ...
11     WRITE(iu, id=id, asynchronous=yes) data
12     ...
13 END SUBROUTINE async_write
```



Asynchronous I/O - 4

An alternative for calling **WAIT** is to periodically call **INQUIRE** to check the status of the operation and in the meantime keep on doing something else

```
1 LOGICAL :: status
2 ...
3 OPEN(10, ..., asynchronous=yes)
4 WRITE(10,..., id=id, asynchronous=yes) A
5 DO WHILE (!status)
6     CALL do_something(...) ! Not involving A
7     INQUIRE (10, id=id, pending=status)
8 END DO
```



Exercises

Edit the file *exampleAsyncIO.F90* to:

- Write the "regular version" (without using `asynchronize I/O`) of the routine *doingAsyncIO*. You can call the new routine *notDoingAsyncIO*.
- Time the four calls to *validWrite*, *invalidWrite*, *doingAsyncIO*, and *notDoingAsyncIO*.

